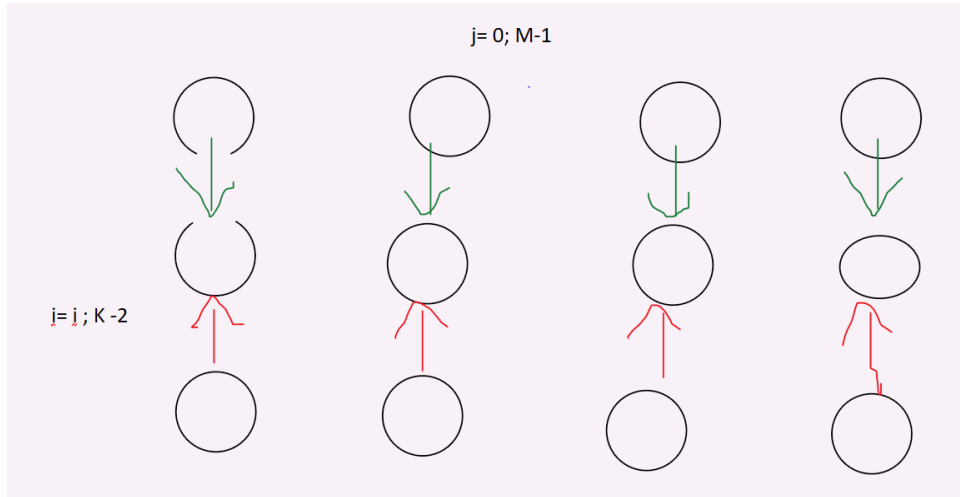


## Problem 1:

- i.
- S1 flow dependent to S3, its ok, in loop
  - S2 flow dependent to S3, its ok, in loop
  - S1 flow dependent to S3
  - S2 anti dependent to S3



- ii. Yes.  
Because All of the "i" iterations are able to start independently from each other.

```
# pragma omp parallel for default (none) private(i, M , a1, a2, j,A) shared(K)
for (int i = 1; i < K-1; i++)
    for (int j = 0; j < M; j++)
        {
            int a1 = A[i-1][j];
            int a2 = A[i+1][j];
            A[i][j] = (a1+a2)/2;
        }
```

## Problem2: Monte Carlo PI Approximation

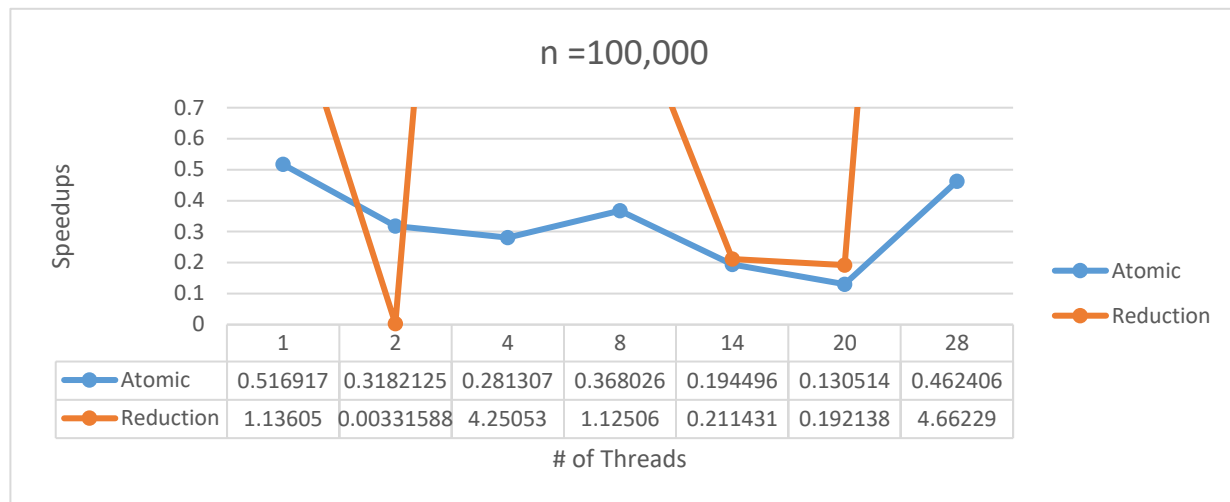


Figure 1

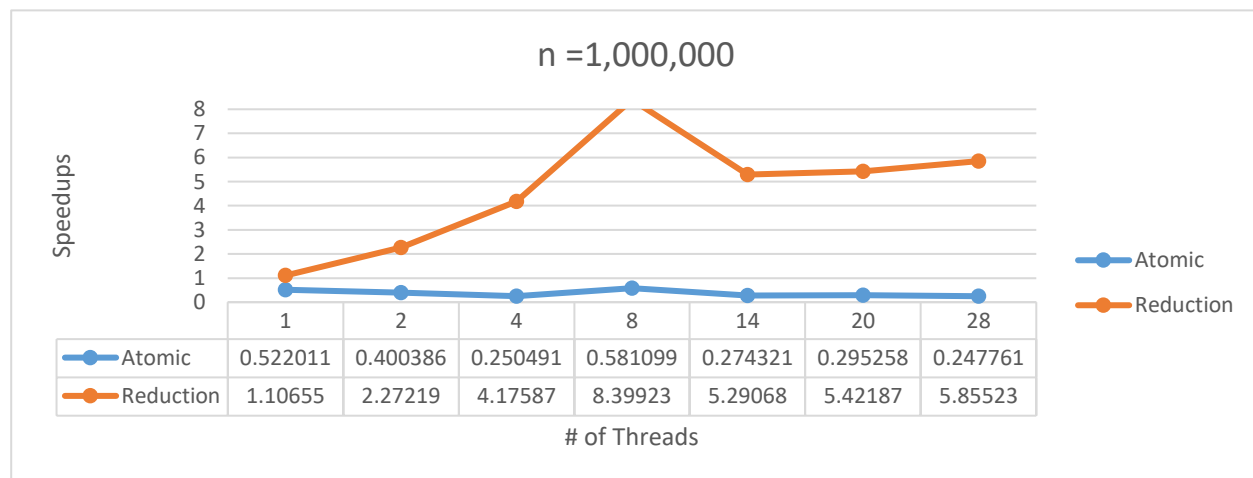


Figure 2

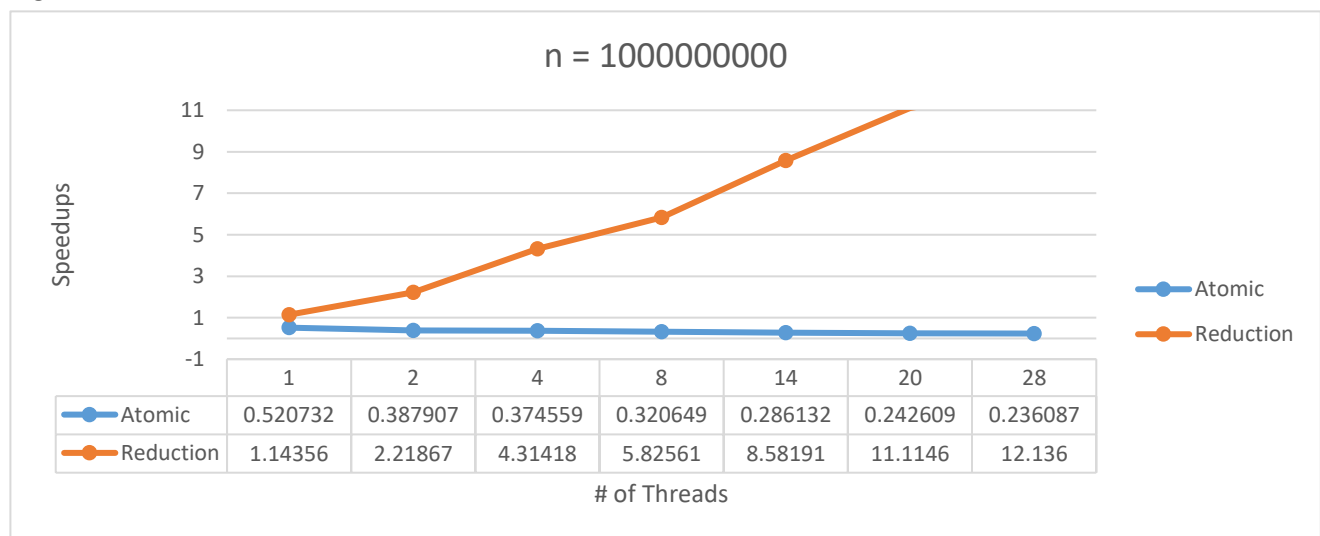


Figure 3

ii.) For the first n size, shown in figure 1, the relative speedups to the sequential code are within the 0 and 1 seconds. On the chart, it can be shown that the reduction directive has a more significant speedup, which means it's the more efficient way of parallelizing source code. For n at 1 million, the speed up range from 0 to 8 and the reduction speed up slope is growing linearly to the number of threads where the atomic is staying relatively constant. Finally, for n at 1 billion the significance change is between 0 and 12. It is clear to see that as n grows, so does the efficiency of using the reduction directive in openMP in comparison to the atomic directive. For the random seed, I used the i variable since it is a private variable for each thread. I originally tried to do the seed as the thread id, but it would sometimes result in an invalid computation, where choosing i, always generates valid approximations.

### Problem3 : Count Sort Algorithm

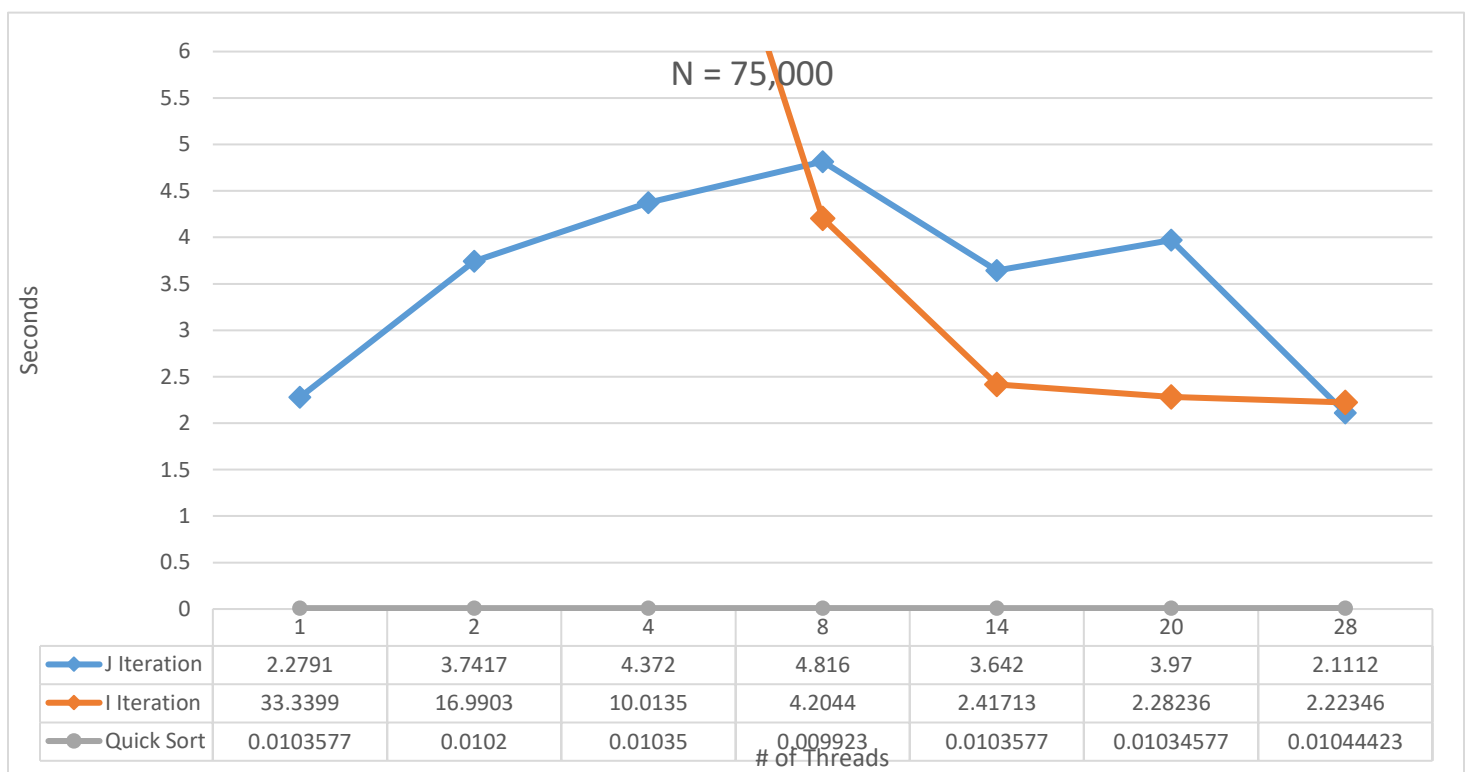


Figure 1

iii.) Both I and J parallelized versions run slower overall in comparison to the sequential quick sort. I believe this makes sense because even if you parallelize an  $n^2$  algorithm, only the constants will change but the overarching asymptotic time complexity will stay the same. After a certain cross over point, as N increases the  $n \log n$  algorithm will continue to run quicker than the  $n^2$  algorithm. Here at 75,000 elements, the speed up is nonexistent in the parallelized code and the j loop runs slower than the i loop linearly until we reach 28 threads, where the i loop speed down is less significant. I'm not exactly sure why. As for the parallelized memcopy, I just created a for loop, and set the contents of an assigned to the contents of temp.

## Bonus Problem: Weight Balancing

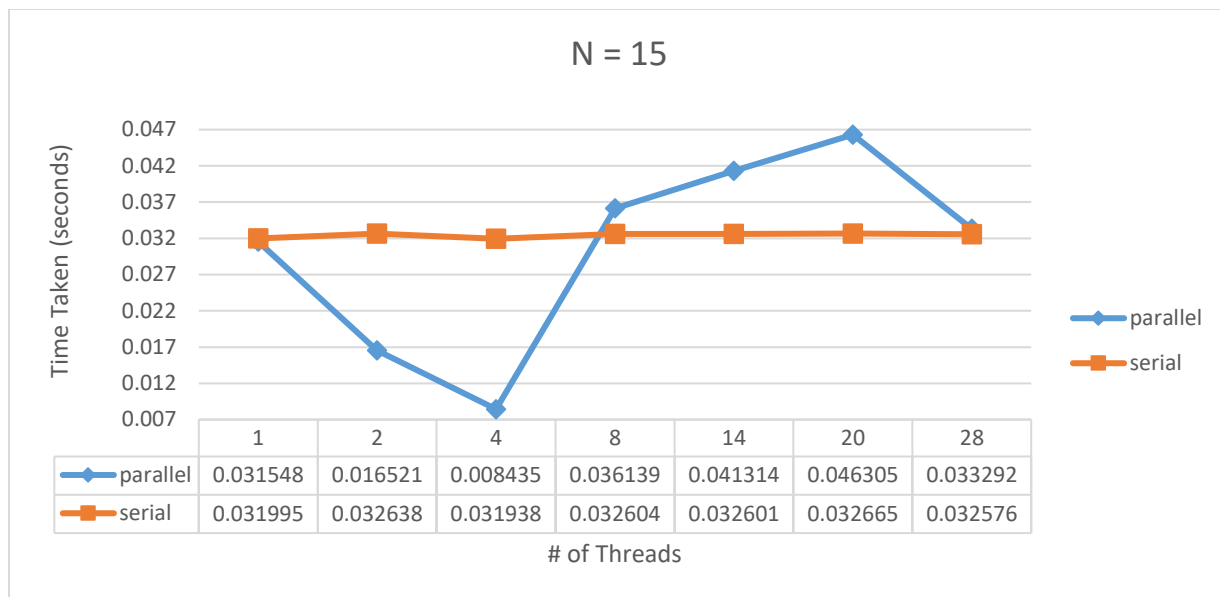


Figure 1

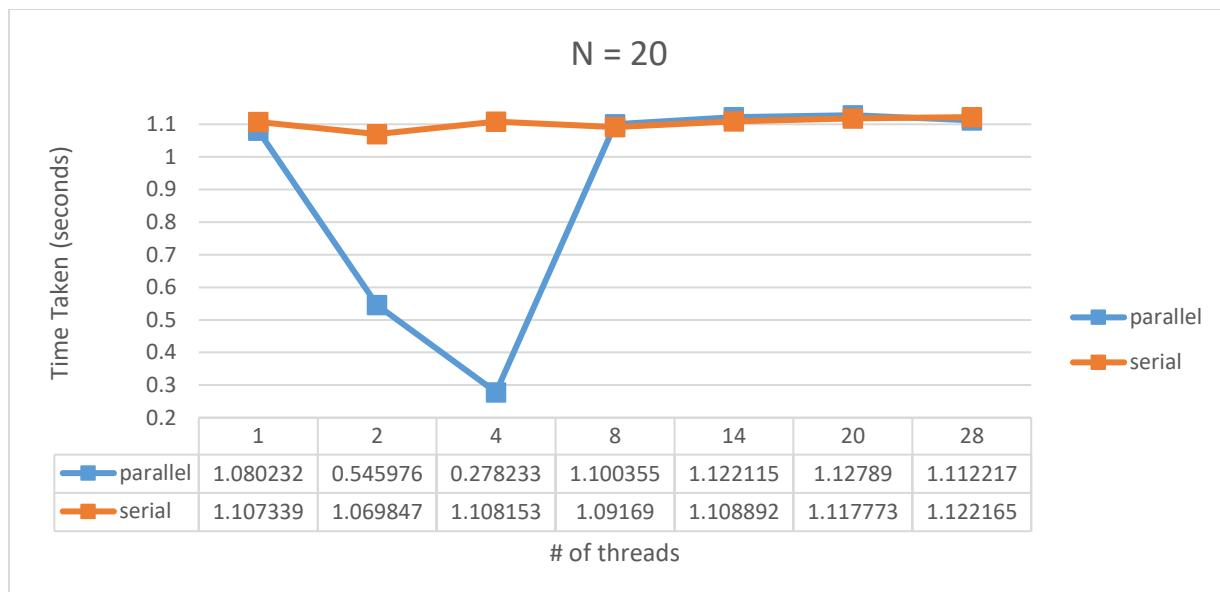


Figure 2

For this problem, I originally believed that the way of not generating too many tasks that created the parallelization code slower in all cases. Something that I decided to do is once my global min has moved to a small weight, I called the serial code, because the overhead of paralleling was no longer needed.

With a consistent N value, thread's with around 2 and 4 have noticeably large speedups. One thing I decided to do, is to limit the number of tasks that can run in parallel. I set that to an arbitrary value of 5. That speed up around thread's close to 5 are greatly significant, with the if clause in the task directive. I

also used the mergeable directive to try to merge any tasks to relieve work for the individual executions. All other thread number arguments, tend to have a small speed down, or practically same speed.

I'm still kind of confused of how that determines such optimal results. In both figures, 4 threads were the number that created such optimal results. I think it should do with how I implemented the if clause. I used a qsub script to grab results, so I believed it wasn't noisy data, however I can't explain what caused figure 1's larger time for 8 thread's spike, and what made figure 2's threads to follow consistently with the serial code.